

KVarN: Variance-Normalized KV-Cache Quantization Mitigates Error Accumulation in Reasoning Tasks

Lorenz K. Muller, Philippe Bich, Chiara Boretti, Hyun-Min Chang, Jiawei Zhuang, Lukas Cavigelli | HUAWEI Computing Systems Lab
<https://arxiv.org/abs/2606.03458>

Research Questions

- Could not extract research questions — LLM analysis was unavailable.

Core Idea

LLM analysis was unavailable.

Limitations

- Analysis unavailable — cannot determine solution limitations.