

SparDA: Sparse Decoupled Attention for Efficient Long-Context LLM Inference

Yaosheng Fu, Guangxuan Xiao, Xin Dong, Song Han, Oreste Villa | NVIDIA

<https://arxiv.org/abs/2606.04511>

Research Questions

- Could not extract research questions — LLM analysis was unavailable.

Core Idea

LLM analysis was unavailable.

Limitations

- Analysis unavailable — cannot determine solution limitations.