

Zone of Proximal Policy Optimization: Teacher in Prompts, Not Gradients

Byung-Kwan Lee, Ximing Lu, Shizhe Diao, Minki Kang, Saurav Muralidharan, Karan Sapra et al. | NVIDIA

<https://arxiv.org/abs/2606.18216>

Research Questions

- Could not extract research questions — LLM analysis was unavailable.

Core Idea

LLM analysis was unavailable.

Limitations

- Analysis unavailable — cannot determine solution limitations.